

# Process Synchronization

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization in Solaris 2 & Windows 2000

# Process cooperation and synchronization

## ■ Process Synchronization

- ...mechanisms to ensure the orderly execution of cooperating processes that share a logical address space, so that data consistency is maintained.

## ■ Why do process cooperate ?


- Modularity: breaking up a system into several subsystems
  - ✓ E.g, an interrupt handler and device driver that need to communicate.
- Convenience: users might want to have several processes to share data
- Speed up: a single program is run as several sub-programs

## ■ How do processes co-operate ?

- Communication abstraction: producers and consumers
  - ✓ Producers produce a piece of information
  - ✓ Customers use this information.

# Background

- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
  - Shared-memory solution to bounded-buffer problem allows at most  $n - 1$  items in buffer at the same time. A solution, where all  $N$  buffers are used is not simple.
    - ✓ Suppose that we modify the producer-consumer code by adding a variable *counter*, initialized to 0 and incremented each time a new item is added to the buffer



# Bounded-Buffer

## ■ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```

# Bounded-Buffer

## ■ Producer process

```
item nextProduced;
```

```
while (1) {  
  while (counter == BUFFER_SIZE)  
    ; /* do nothing */  
  buffer[in] = nextProduced;  
  in = (in + 1) % BUFFER_SIZE;  
  counter++;  
}
```

# Bounded-Buffer

## ■ Consumer process

```
item nextConsumed;
```

```
while (1) {  
  while (counter == 0)  
    ; /* do nothing */  
  nextConsumed = buffer[out];  
  out = (out + 1) % BUFFER_SIZE;  
  counter--;  
}
```

# Bounded Buffer

- Although, both the producer and consumer routines are correct separately they may not function correctly when executed concurrently.

- The statements

**counter++;**  
**counter--;**

must be performed *atomically*.

- Atomic operation means an operation that completes in its entirety without interruption.



# Bounded Buffer

- The statement “**count++**” may be implemented in machine language as:

**register1 = counter**

**register1 = register1 + 1**

**counter = register1**

- The statement “**count—**” may be implemented as:

**register2 = counter**

**register2 = register2 – 1**

**counter = register2**

# Bounded Buffer

- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.

# Bounded Buffer

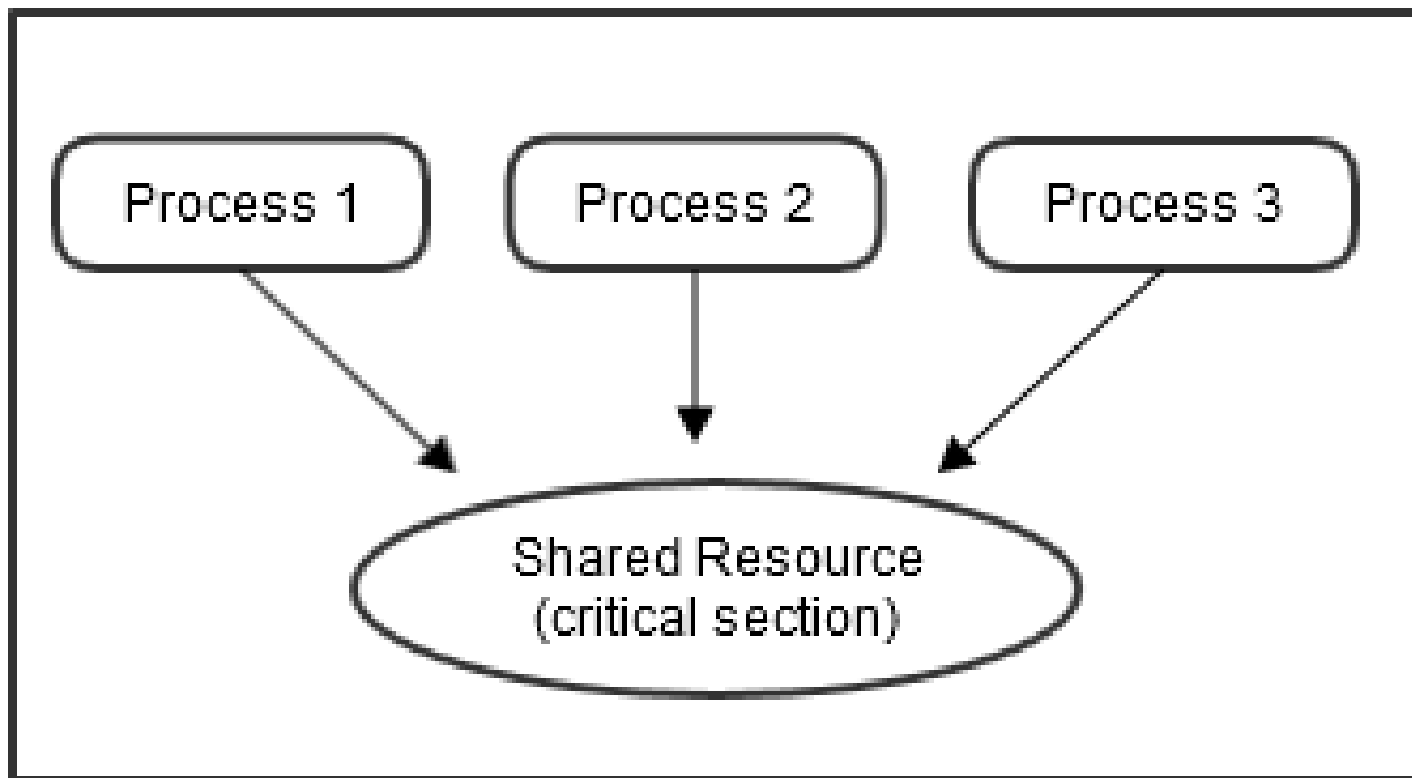
- Assume **counter** is initially 5. One interleaving of statements is:

producer: **register1** = **counter** (*register1* = 5)  
producer: **register1** = **register1** + 1 (*register1* = 6)  
consumer: **register2** = **counter** (*register2* = 5)  
consumer: **register2** = **register2** - 1 (*register2* = 4)  
producer: **counter** = **register1** (*counter* = 6)  
consumer: **counter** = **register2** (*counter* = 4)

- The value of **count** may be either 4 or 6, where the correct result should be 5.

# Race Condition

- **Race condition:** The situation where several processes access – and manipulate shared data concurrently. The final value of the shared data depends upon which process finishes last.
- To prevent race conditions, concurrent processes must be **synchronized**.



# The Critical-Section Problem

- $n$  processes all competing to use some shared data
- Each process has a code segment, called *critical section*, in which the shared data is accessed.
- Problem – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.

# Solution to Critical-Section Problem

- A solution to critical section problem must satisfy the following conditions.
  - **Mutual Exclusion.** If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical section.
  - **Progress.** At least one process requesting entry into CS will be able to enter it if there is no other process in it..
  - **Bounded Waiting.** No process waits indefinitely to enter CS once it has requested entry.
- Assume that each process executes at a nonzero speed
- No assumption concerning relative speed of the  $n$  processes.

# Two approaches

- Several kernel-level processes may active at a time
  - Example: Data structure “List of open files”
- Kernel developers should ensure that OS is free from race conditions.
- Two approaches are ued
- Non-preemptive kernel
  - A non-preemptive kernel does not allow a process running in the kernel mode to be preempted.
  - Kernel mode process runs until it exists kernel mode, blocks, or voluntarily yields the control of CPU
  - Free from race conditions
- Preemptive kernel
  - A preemptive kernel allows a process to be pre-empted while it is running in kernel mode.
  - Should be carefully designed
  - Difficult to design especially in SMP
- Why we prefer preemptive kernels ?
  - Suitable for realtime programming
  - More responsive as kernel mode process can not run for a longer time.
- WINDOWS XP, WINDOWS 2000, Prior to LINUX 2.6 are non-preemptive
- Solaris and IRIX are preemptive



## Mutual exclusion: Software approaches

- Software approaches can be implemented
- Assume elementary mutual exclusion at the memory access level.
  - Simultaneous access to the same location in main memory are serialized in some order.
- Beyond this, no other support in the hardware, OS, programming language is assumed.


# Two process solution

## Initial Attempts to Solve Problem

### Dekker's algorithm

- Reported by Dijkstra, 1965.
- Only 2 processes,  $P_0$  and  $P_1$
- General structure of process  $P_i$  (other process  $P_j$ )

```
do {  
    entry section  
    critical section  
    exit section  
    reminder section  
} while (1);
```
- Processes may share some common variables to synchronize their actions.



# Algorithm 1

- Shared variables:

- **int turn;**  
initially **turn = 0**

- Turn variable

**P0**

while (turn != 0) ;

*/\* Do nothing \*/*

*critical section*

turn = 1;

remainder section

**P1**

while (turn != 1);

*/\* Do nothing \*/*

*critical section*

turn = 0;

remainder section

- Shared variable *turn* indicates who is allowed to enter next, can enter if *turn = me*
  - On exit, point variable to other process
  - Deadlock if other process never enters

- +Satisfies mutual exclusion: Only one process can enter in CS

- -It does not satisfy the progress requirement, as it requires strict alternation of processes to enter CS.

- The pace of execution is dictated by slower process.

- If turn=0, P1 is ready to enter into CS, P1 can not do so, even though P0 may be in the RS.

- If one process fails in CS or RS, other process is blocked permanently.

# Algorithm 2

## ■ Problem with Alg1

- It does not retain sufficient information about the state of each process.
- Alg1 remembers only which process is allowed to enter the CS.

## ■ To solve this problem, variable turn is replaced by **boolean flag[2]**; flag[0] is for P0; and flag[1] is for P1.

## ■ Each process may examine the other's flag but may not alter it.

## ■ When a process wishes to enter CS, it periodically checks other's flag until that flag is false (other process is not in CS)

## ■ The process sets its own flag true and enters CS.

## ■ When it leaves CS, it sets its flag to false.

## Algorithm 2...

- initially **flag [0] = flag [1] = false.**

■ **P0**

```
while ( flag[1] ) ;
```

```
/* Do nothing */
```

```
flag[0] = true;
```

```
critical section
```

```
flag[0] = false;
```

**P1**

```
while ( flag[0] )
```

```
/* Do nothing */
```

```
flag[1] = true;
```

```
critical section
```

```
flag[1] = false;
```

- Mutual exclusion is satisfied.
- If one process fails outside CS the other process is not blocked.
- Sometimes, the solution is worst than previous solution.
  - It does not even **guarantee ME.**
    - ✓ P0 executes the **while** statement and fins flag[1] set to false.
    - ✓ P1 executes the **while** statement and fins flag[0] set to false.
    - ✓ P0 sets flag[0] to true and enters its CS.
    - ✓ P1 sets flag[1] to true and enters its CS.

# Algorithm 3

- Interchange the first two statements.

- Busy Flag Modified

**P0**

flag[0] = true;

while ( flag[1] );

*/\* Do nothing \*/*

*critical section*

flag[0] = false;

**P1**

flag[1] = true;

while ( flag[0] );

*/\* Do nothing \*/*

*critical section*

flag[1] = false;

- Guarantees ME

- Both processes set their flags to true before either has executed the **while** statement, then each will think the other has entered CS causing deadlock.

# Correct solution (1)

- Combining the key ideas of previous algorithms
- Dekker's Algorithm
  - Use *flags* for mutual exclusion, *turn* variable to break deadlock
  - Handles mutual exclusion, deadlock, and starvation
- Dekker's Algorithm
- Initial state: `flag[0]=flag[1]=false; turn=1`

**P0**

```
flag[0] = true;
```

```
while ( flag[1] )  
    if (turn==1)  
    {
```

```
        flag[0]=false;  
        while (turn==1)  
            /* do nothing */  
        flag[0]=true;
```

```
    }
```

```
/* critical section */
```

```
turn=1;
```

```
flag[0] = false;
```

```
remainder section
```

**P1**

```
flag[1] = true;
```

```
while ( flag[0] )  
    if (turn==0)  
    {
```

```
        flag[1]=false;  
        while (turn==0)  
            /* do nothing */  
        flag[1]=true;
```

```
    }
```

```
/* critical section */
```

```
turn=0;
```

```
flag[1] = false;
```

```
remainder section
```



## Correct solution (2)

### ■ Peterson's Algorithm

### ■ Initial state: flag[0]=flag[1]=false;

**P0**

flag[0] = true;

turn = 1;

while ( flag[1] && turn==1)

    /\* Do Nothing \*/;

*critical section*

flag[0] = false;

*remainder section*

**P1**

flag[1] = true;

turn = 0;

while ( flag[0] && turn==0)

    /\* Do nothing \*/;

*critical section*

flag[1] = false;

*remainder section*

# Correct solution

- We need to show that
  - ME is preserved
  - The progress requirement is satisfied
  - The bounded-waiting requirement is met.
- **ME is preserved**
  - If both processes enter the CS both  $\text{flag}[0] == \text{flag}[1] == \text{true}$
  - Both could not execute while loop successfully as turn is either 0 or 1.
- **Progress.**
  - While P1 exits CS it sets  $\text{flag}[1] = \text{false}$ , allowing P0 to enter CS.
  - P1 and P0 will enter the CS (Progress)
- **Bounded waiting:** P1 will enter the CS after at most one entry by P0 and vice versa.

# Multi-process solution: Bakery Algorithm

Critical section for  $n$  processes

- Based on scheduling algorithm commonly used in bakeries.
  - On entering the store the customer receives the number.
  - The customer with the lowest number is served.
  - Customers may receive the same number, then the process with the lowest name is served first.
- Before entering its critical section, process receives a number. Holder of the smallest number enters the critical section.
- If processes  $P_i$  and  $P_j$  receive the same number, if  $i < j$ , then  $P_i$  is served first; else  $P_j$  is served first.
- The numbering scheme always generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5...

# Bakery Algorithm

- var: choosing: array[0...n-1] of boolean.
- Notation  $<\equiv$  lexicographical order (ticket #, process id #)
  - $(a,b) < c,d$  if  $a < c$  or if  $a = c$  and  $b < d$
  - $\max(a_0, \dots, a_{n-1})$  is a number,  $k$ , such that  $k \geq a_i$  for  $i=0, \dots, n-1$
- Shared data
  - boolean choosing[n];**
  - int number[n];**
  - Data structures are initialized to **false** and **0** respectively

# Bakery Algorithm

```
do {  
    choosing[i] = true;  
    number[i] = max(number[0], number[1], ..., number [n – 1])+1;  
    choosing[i] = false;  
    for (j = 0; j < n; j++) {  
        while (choosing[j]) ;  
        while ((number[j] != 0) && (number[j,j] < number[i,i])) ;  
    }  
    critical section  
    number[i] = 0;  
    remainder section  
} while (1);
```

- Consider  $P_i$  in its CS and  $P_k$  is trying to enter CS
- When  $P_k$  enters second while statement for  $j=i$ , it finds that
  - $\text{number}[i] \neq 0$
  - $(\text{number}[i], i) < (\text{number}[k], k)$
  - So it waits until  $P_i$  leaves CS
- FCFS is followed.

# Mutual exclusion: hardware solution

- In the uni-processor system, it is sufficient to prevent a process from being interrupted.

```
while (true){  
    /* disable interrupts */  
    /* Critical section */  
    /* enable interrupts */  
    /* remainder */  
}
```

- Since CS can not be interrupted ME is guaranteed.
- The efficiency decreases
- It can not work in multi-processor environments
  - More than one process is executing at a time.

# Special machine instructions

- In multi-processor configuration, several processes share access to a common main memory.
- At the hardware level, access to a memory location excludes any other access to that same memory location.
- Processor designers have proposed several machine instructions to carry out two actions atomically (single cycle).
  - **Reading and writing**
  - **swapping**

# Test and set instruction

- Test and modify the content of a word atomically

```
boolean testset (int i)
```

```
{  
  if (i==0)  
  {  
    i=1;  
    return true;  
  }  
  else  
  {  
    return false  
  }  
}
```

- This instruction sets the value of 'i', if the value=0 and returns true. Otherwise the value is not changed and false is returned.



# Mutual Exclusion with Test-and-Set

- Shared data:

```
boolean lock = false;
```

- `void P(int i)`

```
do {
```

```
    while (TestAndSet(lock)==false)
```

```
        /* do nothing*/;
```

```
        critical section
```

```
    lock = false;
```

```
    remainder section
```

```
}
```

```
void main()
```

```
{
```

```
    lock=false;
```

```
    parbegin(P1(), P(2),...,P(n));
```

```
}
```

# Test-and-Set: Correctness

## ■ Mutual exclusion

- A shared variable lock is set to false
- The only process  $P_i$  that enters CS that finds lock as false and sets it to true.
- All other processes trying to enter CS so into a busy waiting mode and finds lock as false.
- When process leaves C it resets lock to false.
- When  $P_i$  exits lock is set to false so the next process  $P_j$  to execute instruction find test-and-set=false and will enter the CS.

## ■ Progress

- Trivially true

## ■ Unbounded waiting

- Possible since depending on the timing of evaluating the test-and-set primitive.
- Does not guarantee fairness.

# Swap instruction

- Atomically swap two variables.

```
void swap(boolean &a, boolean &b) {  
    boolean temp = a;  
    a = b;  
    b = temp;  
}
```

# Mutual Exclusion with Swap

- Shared data (initialized to **false**):  
    **boolean lock;**  
    **boolean waiting[n];**
- Process  $P_i$   
    **do {**  
        **key = true;**  
        **while (key == true)**  
            **Swap(lock,key);**  
        critical section  
        **lock = false;**  
        remainder section  
    **}**

# SWAP: Correctness

- **Similar to Test-and-set**
- **Mutual exclusion**
- **Progress**
  - **Trivially true**
- **Unbounded waiting**
  - **Possible since depending on the timing of evaluating the test-and-set primitive.**
  - **Does not guarantee fairness.**

# Can we get bounded waiting ?

- Introduce a boolean array called waiting of size n and boolean variable key

## ■ Entry

- waiting[i]:=true;
- key:=true;
- while (waiting[i] and key ) do
  - ✓ key := test-and-set(lock);
- waiting[i]:=false;
- execute CRITICAL SECTION

## ■ Exit

- Find the next process j that has waiting[j]=1 stepping through waiting.
- Set waiting[j]:=false;
- Process  $P_j$  immediately enter the CS.
- If no process exists, set lock=false;

# Can we get bounded waiting ?....

- Every (interested)  $P_i$  executes that test&set at least once.
- $P_i$  enters the critical section provided:
  - Key is false in which case there is no process in CS.
- Or
  - If it was waiting, because  $\text{waiting}[i]$  was reset to false by the unique process that was blocking it in the critical section.
  - Either of the above events occur exactly once and hence mutual exclusion.

# Properties of machine instruction approach

■ +ve


- Any number of processes
- Simple and easy
- Can support multiple CSs.

■ -ve

□ **Busy waiting** is employed

- ✓ The process is waiting and consuming processor time.
- Starvation is possible.
  - ✓ The selection of waiting process is arbitrary.
- Deadlock is possible due to priority
  - ✓ P1 enters CS and interrupted by higher priority process P2 which is trying to enter CS.
  - ✓ P2 can not get CS unless P1 is out and P1 can not be dispatched due to low priority.





# Mutex Locks

- Test-set locks are also called mutex llocks
- Mutex means mutual exclusion.

# Semaphores: Dijkstra; 1965

- Two and more processes can cooperate by means of simple signals, such that a process is forced to stop at a specified place until it has received a specific signal.
- For signaling, special variables called semaphores are used
- A semaphore is a synchronization tool.
- A semaphore is an integer variable that is accessed only through two standard atomic operations: **wait and signal**.
- To transmit a signal to semaphore S, a process executes the primitive *signal(S)* primitive.
- To receive a signal via semaphore S, the process executes *wait(S)* primitive.

# Semaphores: Dijkstra 1965

## Classical or first definition

- A semaphore is initialized to a non-negative value
- The **wait** operation decrements the semaphore value. If the integer value is negative the process waits.
- The **signal** operation increments the semaphore value. If the value is not positive, then process which is blocked by a wait operation is gets the access to CS.
- **The wait and signal are assumed to be atomic.**
- Semaphore  $S$  – integer variable
- can only be accessed via two indivisible (atomic) operations

*wait* ( $S$ ):

**while  $S \leq 0$  do *no-op*;**  
 **$S--$ ;**

*signal* ( $S$ ):

**$S++$ ;**

# Critical Section of $n$ Processes

- Shared data:

**semaphore mutex;** // initially *mutex* = 1

- Process  $P_i$ :

```
do {  
    wait(mutex);  
    critical section  
    signal(mutex);  
    remainder section  
} while (1);
```

- Modifications to the integer value of the semaphore in the wait and signal operations must be executed indivisibly.

# Semaphore Implementation

- The classical definition requires busy waiting.
- While a process is in CS, the other process must loop continuously in the entry code.
- Busy waiting wastes CPU cycles.
- This type of semaphore is called **spinlock**: process spins while waiting for a lock.
  - Advantage of spinlock: no context switch
  - When locks are expected to be held for short times, spinlocks are useful.
- To overcome the need for busy waiting, we can modify the definition of the wait and signal semaphore operations.
- If a process executes wait operation and finds the semaphore operation is not positive, it must wait.
  - Rather than busy waiting it must **block** itself.
  - The **block** operation puts the process into waiting queue of semaphore and process is switched to waiting state.
- A process that is blocked waiting on a semaphore S, should be restarted when some other process executes signal operation.
- The process is restarted with **wakeup** operation.

# Semaphore Implementation

- Define a semaphore as a record

```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore;
```

- Assume two simple operations:
  - **block** suspends the process that invokes it.
  - **wakeup(*P*)** resumes the execution of a blocked process *P*.

# Implementation

- Semaphore operations now defined as

*wait(S):*

```
S.value--;  
if (S.value < 0) {  
    add this process to S.L;  
    block;  
}
```

*signal(S):*

```
S.value++;  
if (S.value <= 0) {  
    remove a process P from S.L;  
    wakeup(P);  
}
```

- Wait and signal operations are system calls.



# Semaphore as a General Synchronization Tool

- Execute  $B$  in  $P_j$  only after  $A$  executed in  $P_i$
- Use semaphore  $flag$  initialized to 0
- Code:

$P_i$	$P_j$
$\vdots$	$\vdots$
$A$	$wait(flag)$
$signal(flag)$	$B$

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.

- Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$	$P_1$
<i>wait(S);</i>	<i>wait(Q);</i>
<i>wait(Q);</i>	<i>wait(S);</i>
$\vdots$	$\vdots$
<i>signal(S);</i>	<i>signal(Q);</i>
<i>signal(Q);</i>	<i>signal(S);</i>

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

# Two Types of Semaphores

- *Counting* semaphore – integer value can range over an unrestricted domain.
- *Binary* semaphore – integer value can range only between 0 and 1; can be simpler to implement.

# Binary Semaphores

- A binary semaphore is a semaphore with an integer value that can range only between 0 and 1

- It is simple to implement.

- Type binary semaphore = **record**  
    value:(0,1)  
    queue: list of processes

**end;**

- var s: binary semaphore

- **waitB(s):**

**if** s.value=1 **then**

        s.value=0

**else**

**begin**

            place this process in s.queue;

            block this process;

**end;**

- **signalB(s):**

**if** s.queue is empty **then**

        s.value=1

**else**

**begin**

            remove (wakeup) the process from  
            s.queue;

            place this process in the ready list.

**end;**

# Implementing *S* as a Binary Semaphore

- Can implement a counting semaphore *S* as a binary semaphore.
- Data structures:  
**binary-semaphore S1, S2;**  
**int C;**
- Initialization:  
**S1 = 1**  
**S2 = 0**  
**C = initial value of semaphore S**

```
□ wait operation
  wait(S1);
  C--;
  if (C < 0) {
    signal(S1);
    wait(S2);
  }
  signal(S1);
```

```
□ signal operation
  wait(S1);
  C ++;
  if (C <= 0)
    signal(S2);
  else
    signal(S1);
```

## *Counting semaphores*

*wait(S):*

```
  S.value--;
  if (S.value < 0) {
    add this process to
    S.L;
    block;
  }
```

*signal(S):*

```
  S.value++;
  if (S.value <= 0) {
    remove a process P
    from S.L;
    wakeup(P);
  }
```

# Classical Problems of Synchronization

- Bounded-Buffer Problem

- Readers and Writers Problem

- Dining-Philosophers Problem

# Bounded-Buffer Problem

- Used to to illustrate the power of synchronization techniques
- We assume that the buffer consists of  $n$  buffers, each capable of holding an item.
- The mutex semaphore provides mutual exclusion access to buffer which is initialized to the value 1.
- The **empty** and **full** semaphores count the number of empty and full buffers which are initialized to  $n$  and zero respectively.
- Shared data  
**semaphore full, empty, mutex;**
- Initially:  
**full = 0, empty =  $n$ , mutex = 1**

# Bounded-Buffer Problem Producer Process

```
do {  
    ...  
    produce an item in nextp  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (1);
```



# Bounded-Buffer Problem Consumer Process

```
do {  
    wait(full)  
    wait(mutex);  
    ...  
    remove an item from buffer to nextc  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```

- Producer is producing full buffers for the consumer and consumer is producing empty buffers for the consumer.

# Readers-Writers Problem

- Problem: A data object (file or record) is shared among several concurrent processes.
  - Some want to read and others want to update it.
- **Readers:** processes interested in reading.
- **Writers:** processes interested in writing.
- Two readers can access shared data object simultaneously.
- But a writer and reader can access shared data object simultaneously
  - problems may occur!
- To protect from these problems, writers should have an exclusive access to the shared object.
- This synchronization problem is referred to as readers-writers problem.
- It is a different kind of synchronization problem.
- The readers-writers problem has several variations.
  - Simple one: No reader will be kept waiting unless writer has obtained permission to write.

# Readers-Writers Problem

- Semaphores used: **mutex** and **wrt**
- The semaphore **wrt** is common to reader and writer.
- Semaphore **mutex** is used to update **readcount**.
- **readcount** keeps track of how many are reading the object.
- Shared data

**semaphore mutex, wrt;**

Initially

**mutex = 1, wrt = 1, readcount = 0**

# Readers-Writers Problem Writer Process

**wait(wrt);**

...

writing is performed

...

**signal(wrt);**

# Readers-Writers Problem Reader Process

```
wait(mutex);  
readcount++;  
if (readcount == 1)  
    wait(wrt);  
signal(mutex);  
...  
reading is performed  
...  
wait(mutex);  
readcount--;  
if (readcount == 0)  
    signal(wrt);  
signal(mutex);
```


- Writers can be starved if there is a continuous sequence of readers.

# Readers-Writers Problem

- Can the producer/consumer problem be considered as a case of the readers/writers problem with a writer as a producer and reader as a consumer ?
- The answer is no
- The producer is not just a writer
  - It must read queue of pointers to determine where to write the next item and it must determine if the buffer is full.
- Similarly the consumer is not a reader
  - It must adjust queue pointers to show that it has removed a unit from the buffer.

# Dining-Philosophers Problem

- Five philosophers spend their lives on thinking and eating.
- They share a common circular table surrounded by five chairs.
- Five single chopsticks are available.
- Whenever a philosopher wants to eat, he tries to pick up two chopsticks that are closest to him/her.
- A philosopher can not pick the chopstick in the hand of neighbor.
- After finishing, the philosopher puts back the chopsticks and starts thinking.
- It is simple representation of the need to allocate several resources among several processes in a **deadlock and starvation free manner**.



# Dining-Philosophers Problem

- Shared data

```
semaphore chopstick[5];  
Initially all values are 1
```

- Philosopher  $i$ :

```
do {  
    wait(chopstick[i])  
    wait(chopstick[(i+1) % 5])  
    ...  
    eat  
    ...  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    ...  
    think  
    ...  
} while (1);
```

- The solution creates a deadlock



# Barbershop Problem

- 3 barbers, each with a barber chair
  - Haircuts may take varying amounts of time
- Sofa can hold 4 customers, max of 20 in shop
- Customers wait outside if necessary
- When a chair is empty:
  - Customer sitting longest on sofa is served
  - Customer standing the longest sits down
- After haircut, go to cashier for payment
  - Only one cash register
  - Algorithm has a separate cashier, but often barbers also take payment
    - ✓ This is also a critical section

# Barbershop Problem

- The main body of the program activates 50 customers, 3 barbers, and the cashier process. Synchronization operators.
  - Shop and sofa capacity: the capacity of shop and the capacity of the sofa are governed by the semaphores **max\_capacity** and **sofa**.
    - ✓ When customer enters max\_capacity decremented by one.
    - ✓ When a customer leaves it is incremented.
    - ✓ Wait and signal operations are surround the actions of sitting and getting\_up from sofa.
  - **Barber chair capacity:**
    - ✓ There are three barber chairs; the semaphore barber\_chair assures that no more than three customers attempt to obtain service at a time.
    - ✓ A customer will not get up from the sofa until at least one chair is free.
  - **Ensuring customers are in the barber chair:** The semaphore cust\_ready provides a wakeup signal for a sleeping barber indicating that the customer has just taken the chair.
  - **Holding customers in barber chair:** once seated the customer remain in the chair until the barber gives the signal that haircut is complete, using the semaphore finished.
  - **Limiting one customer to a barber chair:** the semaphore barber\_chair is intended to limit the number of customers in barber chairs to three. The semaphore leave\_b\_chair is used to synchronize sitting.
  - **Paying and receiving:** payment and receipt semaphores are used to synchronize the operations.
  - **Coordinating barber and cashier functions:** To save money the barber shop does not employ a separate cashier. Each barber is required to perform that task when not cutting hair. The semaphore coord ensures the barbers perform only one task at a time.

# Barbershop Problem

Semaphore	Wait operation	Signal operation
<b>max_capacity</b>	Customer waits for a room to enter shop.	Exiting customer signals customer waiting to enter
<b>sofa</b>	Customer waits for seat on sofa	Customer leaving sofa signals customer waiting for sofa
<b>barber_chair</b>	Customer waits for empty barber chair	Barber signals when that barber's chair is empty
<b>Cust_read</b>	Barber waits until customer is in the chair	Customer signals barber that customer is in the chair
<b>finished</b>	Customer waits until his haircut is complete.	Barber signals when done cutting hair of his customer.
<b>leave_b_chair</b>	Barber waits until customer gets up from the chair	Customer signals barber when customer gets up from chair.
<b>payment</b>	Cashier waits for a customer to pay	Customer signals cashier that he has paid.
<b>receipt</b>	Customer waits for a receipt for a payment	Cashier signals that payment has been accepted.
<b>coord</b>	Wait for a barber resource to be free to be free perform either the hair cutting or cashiering function.	Signal that a barber resource is free.

# Barbershop

```
program      barbershop1;
var          max_capacity: semaphore (:=20);
            sofa: semaphore (:=4);
            barber_chair, coord: semaphore (:=3);
            cust_ready, leave_b_chair, payment, receipt: semaphore (:=0)
```

```
procedure customer;
var custnr: integer;
begin
    wait (max_capacity );
    enter shop;
```

```
    wait( sofa );
    sit on sofa;
    wait( barber_chair );
    get up from sofa;
    signal( sofa );
    sit in barber chair;
    wait( mutex2 );
    signal( cust_ready );
    wait( finished[custnr] );
    leave barber chair;
    signal( leave_b_chair );
    pay;
    signal( payment );
    wait( receipt );
    exit shop;
    signal( max_capacity );
end;
```

```
procedure barber;
var b_cust: integer
begin
    repeat
        wait( cust_ready );
```

```
        wait( coord );
        cut hair;
        signal( coord );
        signal( finsihed[b_cust] );
        wait( leave_b_chair );
        signal( barber_chair );
    forever
end;
```

```
procedure cashier;
begin
    repeat
        wait( payment );
        wait( coord );
        accept payment;
        signal( coord );
        signal( receipt );
    forever
```

```
Void main()
{
    count=0;
    Parbegin {customer... 50 times,...customer,
    Barber, barber,barber, cashier)
}
```

```

/* program barbershop1 */
semaphore max_capacity = 20;
semaphore sofa = 4;
semaphore barber_chair = 3;
semaphore coord = 3;
semaphore cust_ready = 0, finished = 0, leave_b_chair = 0, payment = 0, receipt = 0;

void customer ()
{
    wait(max_capacity);
    enter_shop();
    wait(sofa);
    sit_on_sofa();
    wait(barber_chair);
    get_up_from_sofa();
    signal(sofa);
    sit_in_barber_chair;
    signal(cust_ready);
    wait(finished);
    leave_barber_chair();
    signal(leave_b_chair);
    pay();
    signal(payment);
    wait(receipt);
    exit_shop();
    signal(max_capacity)
}

void barber()
{
    while (true)
    {
        wait(cust_ready);
        wait(coord);
        cut_hair();
        signal(coord);
        signal(finished);
        wait(leave_b_chair);
        signal(barber_chair);
    }
}

void cashier()
{
    while (true)
    {
        wait(payment);
        wait(coord);
        accept_pay();
        signal(coord);
        signal(receipt);
    }
}

void main()
{
    parbegin (customer, . . . 50 times, . . . customer, barber, barber, barber,          cashier);
}

```

Figure 16: An unfair barbershop

Another method to avoid the unfairness is to number the barber chairs so that less semaphores are needed, but how? Think about it!

# Barbershop Problem

- The preceding solution is unfair.
- The customers are served in the order they enter the shop.
- If one barber is very fast and one of the customer is quite bald.
- The problem can be solved with more semaphores.
  - We assign unique customer number `id` to each customer.
  - The semaphore `mutex1` protects access to global variable `count`.
- The semaphore `finished` is refined to be an array of 50 semaphores.
  - Once a customer seated in a barber chair, he executes `wait(finished[custnt])` to wait in his own unique semaphore.
- Please see the solution in William Stallings book (pp 229-234)

# Fair Barbershop

```
program
var
    barbershop2;
    max_capacity: semaphore (:=20);
    sofa: semaphore (:=4);
    barber_chair, coord: semaphore (:=3);
    mutex1, mutex2: semaphore (:=1);
    cust_ready, leave_b_chair, payment, receipt: semaphore (:=0);
    finished: array [1..50] of semaphore (:=0);
    count: integer;
```

```
procedure customer;
var custnr: integer;
begin
    wait (max_capacity );
    enter shop;
    wait( mutex1 );
    count := count + 1;
    custnr := count;
    signal( mutex1 );
    wait( sofa );
    sit on sofa;
    wait( barber_chair );
    get up from sofa;
    signal( sofa );
    sit in barber chair;
    wait( mutex2 );
    enqueue1( custnr );
    signal( cust_ready );
    signal( mutex2 );
    wait( finished[custnr] );
    leave barber chair;
    signal( leave_b_chair );
    pay;
    signal( payment );
    wait( receipt );
    exit shop;
    signal( max_capacity );
end;
```

```
procedure barber;
var b_cust: integer
begin
    repeat
        wait( cust_ready );
        wait( coord );
        wait( mutex2 );
        dequeue1( b_cust );
        signal( mutex2 );
        wait( coord );
        cut hair;
        signal( coord );
        signal( finished[b_cust] );
        wait( leave_b_chair );
        signal( barber_chair );
    forever
end;
```

```
procedure cashier;
begin
    repeat
        wait( payment );
        accept payment;
        signal( coord );
        signal( receipt );
    forever
end;
```

```
Void main()
{
    count=0;
    Parbegin {customer... 50 times,...customer,
    Barber, barber,barber, cashier}
}
```

```

/* program barbershop2 */
semaphore max_capacity = 20;
semaphore sofa = 4;
semaphore barber_chair = 3, coord = 3;
semaphore mutex1 = 1, mutex2 = 1;
semaphore cust_ready = 0, leave_b_chair = 0, payment = 0, receipt = 0;
semaphore finished [50] = {0};
int count;

void customer()
{
    int custnr;
    wait(max_capacity);
    enter_shop();
    wait(mutex1);
    count++;
    custnr = count;
    signal(mutex1);
    wait(sofa);
    sit_on_sofa();
    wait(barber_chair);
    get_up_from_sofa();
    signal(sofa);
    sit_in_barber_chair();
    wait(mutex2);
    enqueue1(custnr);
    signal(cust_ready);
    signal(mutex2);
    wait(finished[custnr]);
    leave_barber_chair();
    signal(leave_b_chair);
    pay();
    signal(payment);
    wait(receipt);
    exit_shop();
    signal(max_capacity)
}

void barber()
{
    int b_cust;
    while (true)
    {
        wait(cust_ready);
        wait(mutex2);
        dequeue1(b_cust);
        signal(mutex2);
        wait(coord);
        cut_hair();
        signal(coord);
        signal(finished[b_cust]);
        wait(leave_b_chair);
        signal(barber_chair);
    }
}

void cashier()
{
    while (true)
    {
        wait(payment);
        wait(coord);
        accept_pay();
        signal(coord);
        signal(receipt);
    }
}

void main()
{
    count := 0;
    parbegin (customer, ... 50 times, ... customer, barber, barber, barber,
              cashier);
}

```

Figure 17: An fair barbershop



# Implementing wait() and signal() in Multi-processor Systems

- Disabling interrupts will not work.
- Spinlock is the solution
  - With this we have moved busy waiting from entry section to critical sections of application programs.

# Implementation of Semaphores

- wait and signal operations are atomic.
- No two processes should execute wait and signal on the same semaphore at the same time
- Good Solution: implement through hardware or firmware.
- Other solutions
  - Ensure that only process manipulates “wait” and “signal” operations.
  - One can use Dekker’s algorithm or Peterson’s algorithm
    - ✓ Substantial processing overhead
  - Use one of the hardware supported schemes
    - ✓ Test and set
    - ✓ disabling interrupts (single processor)

# About busy waiting

- Note: we have not eliminated the busy waiting with `wait()` and `signal()` completely.
  - Moved busy waiting from entry section to critical sections of application programs.
  - Furthermore, we have limited busy waiting to critical sections of `wait()` and `signal()` operations.
  - The wait and signal code is very short the amount of busy waiting involved is short.

# Two possible implementations of Semaphores

```
Wait(s)
{
    while(!testset(s.flag)
        /* do nothing */
        s.count--;
    if (s.count < 0)
    {
        place this process in s.queue;
        block this process (set s.flag to 0)
    }
    else
        s.flag=0;
}
```

```
Signal(s)
{
    while(!testset(s.flag)
        /* do nothing */
        s.count++;
    if (s.count <= 0)
    {
        remove a process P from s.queue;
        Place a process P in the ready list
    }
    s.flag=0;
}
```

**With TestSet Instruction**

7.76

```
Wait(s)
{
    Inhibit interrupts
    s.count--;
    if (s.count < 0)
    {
        place this process in s.queue;
        block this process allow interrupts
    }
    else
        allow interrupts;
}
```

```
Signal(s)
{
    Inhibit interrupts;
    s.count++;
    if (s.count <= 0)
    {
        remove a process P from s.queue;
        Place a process P in the ready list
    }
    allow interrupts;
}
```

**With Interrupts**

# Problem with semaphores

- Incorrect use may result in timing errors
- These errors are difficult to detect as these occur if only particular sequence occurs.
- Missing or reverse order.
- It is difficult to produce correct program using semaphores.
- The wait and signal operations are scattered throughout the program and it is difficult to see overall effect of these operations on the semaphores.

# Problem with semaphores (cont.)

## ■ Problems

- Suppose a process interchanges the order in which wait and signal operations on the semaphore are executed

signal (mutex)

CS

wait (mutex)

- ✓ Several processes may be executing in their CS simultaneously.

- Suppose that a process replaces signal(mutex) with wait(mutex)

signal (mutex)

wait(mutex)

CS

CS

signal (mutex)

wait(mutex)

- ✓ Deadlock will occur.

- Suppose a process omits wait(mutex) or signal(mutex) or both.

- ✓ ME is violated or deadlock occurs.

- A critical region and monitor concept is introduced to address this problem

# High level synchronization constructs

- To deal with the type of errors caused by semaphores, high-level constructs have been introduced.
  - Critical region
  - Monitor
- Assumption
  - Process contains some local data
  - Local data can be accessed by only the sequential program that is encapsulated within the same process.
  - Process can not access the local data of another process.

# Monitors

- It is an high level synchronization construct.
- A monitor is a software module consisting of one or more procedures, an initialization sequence and local data.
- Main characteristics
  1. The local data variables are accessible only by the monitor's procedures and not by any external procedure.
  2. A process enters monitor by entering one of its procedures.
  3. Only one process may be executing (active) in the monitor at any time; any other process that has invoked the monitor is suspended while waiting for the monitor to become available.
- 1&2 → object oriented characteristics.
- By enforcing one procedure at a time, the monitor enforces ME facility.



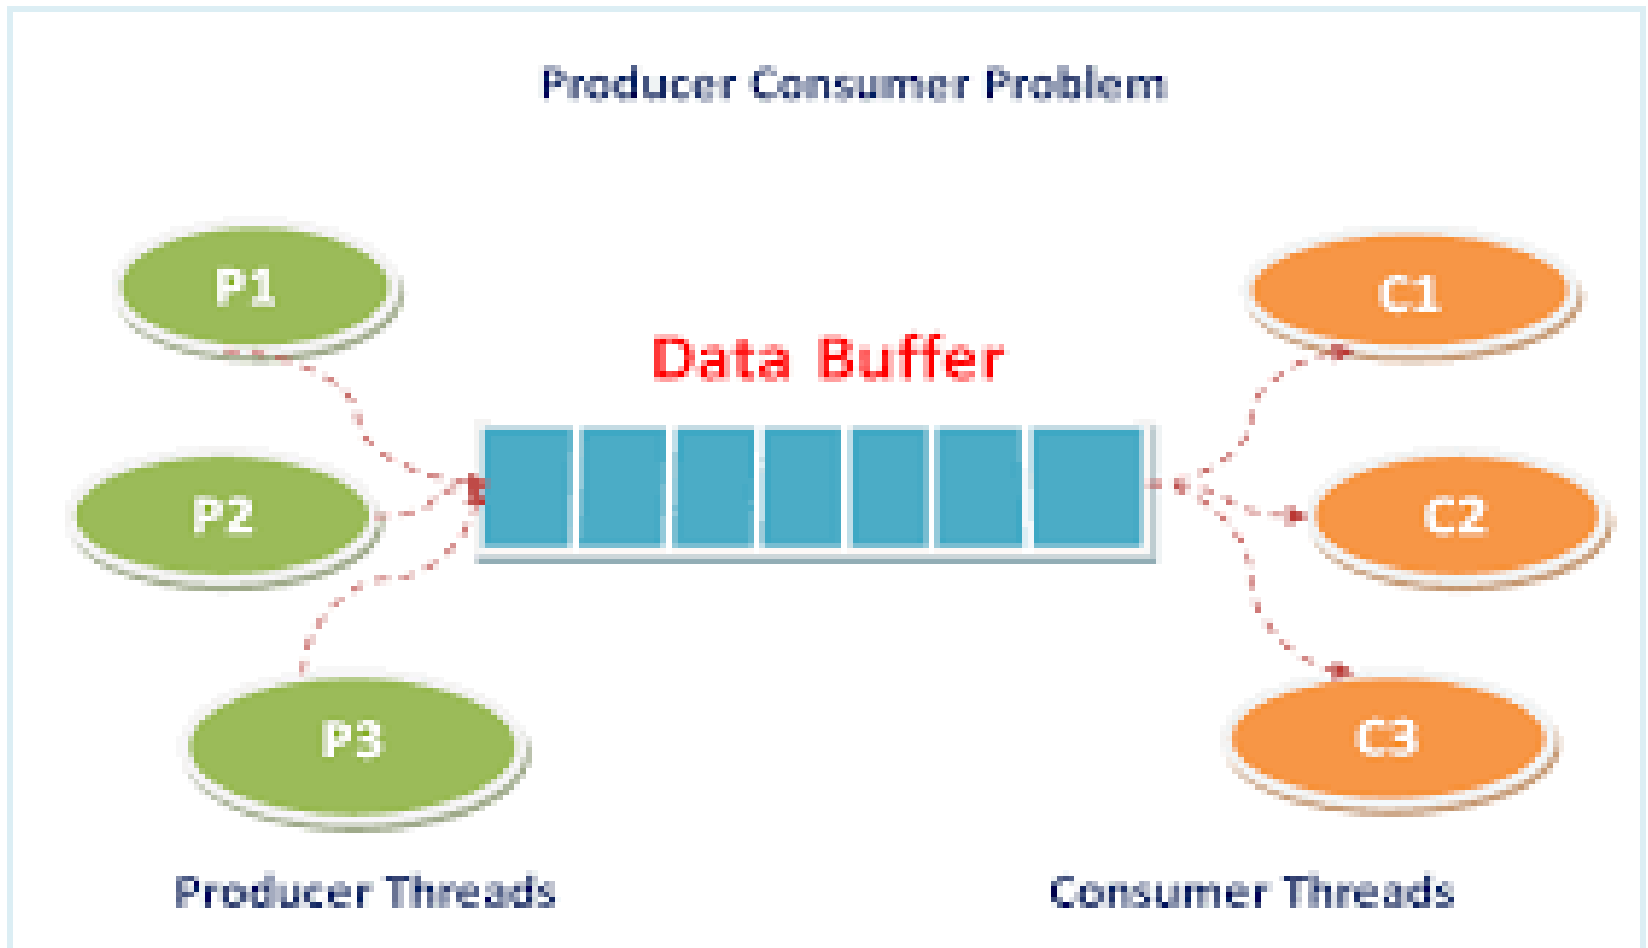
# Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.
- A shared data resource can be protected by placing in the monitor.

```
monitor monitor-name
{
    shared variable declarations
    procedure body P1 (...) {
        . . .
    }
    procedure body P2 (...) {
        . . .
    }
    procedure body Pn (...) {
        . . .
    }
    {
        initialization code
    }
}
```

Two kinds of waiting:

- Mutual exclusion waiting: to avoid race condition
- Conditional waiting: to avoid inconsistency



# Monitors

- To allow a process to wait within the monitor, a **condition** variable must be declared, as

**condition x, y;**

- Condition variable can only be used with the operations **wait** and **signal**.

- The operation


**x.wait();**

means that the process invoking this operation is suspended until another process invokes


**x.signal();**

- The **x.signal** operation resumes exactly one suspended process. If no process is suspended, then the **signal** operation has no effect.

# Schematic View of a Monitor



# Monitor With Condition Variables



# Monitors

- In case of monitors, the monitor construct itself provides ME, but synchronization is provided by the programmer.
- In case of semaphore, both ME and synchronization are provided by the programmer.
- Also , in case of monitors also, it is possible to make mistakes in the synchronization of monitors.
- For example if csignal function is omitted, the processes entering corresponding queue are permanently hung up.
- However, since all synchronization functions are confined to monitor, it is easier to verify the synchronization and detect bugs.
- Once a monitor is correctly programmed, access to the protected resource is correct from all processes.
- With semaphores, resources access is correct only if all of the processes that access the resource are programmed correctly.

# Dining Philosophers Example

## Deadlock free solution

- A philosopher is allowed to pick up his chopsticks only if both of them were available.
- We introduce three states:
  - Enum {thinking, hungry, eating} state[5]
- Philosopher  $i$  can set the variable  $state[i]=\text{eating}$  only if her two neighbors are not eating:  $(state[(i+4)\%5] \neq \text{eating})$  and  $(state[(i+1)\%5] \neq \text{eating})$
- We also declare condition  $self[5];$ 
  - Philosopher  $i$  can delay himself when he is hungry, but unable to obtain the chopsticks he needs.

# Dining Philosophers Example

## Deadlock free solution

```
■ monitor dp
{
    enum {thinking, hungry, eating} state[5];
    condition self[5];
    void pickup(int i)           // following slides
    void putdown(int i) // following slides
    void test(int i)             // following slides

    void init() {
        for (int i = 0; i < 5; i++)
            state[i] = thinking;
    }
}
```



# Dining Philosophers

```
void pickup(int i) {  
    state[i] = hungry;  
    test[i]; // if left and right of i are not eating, then eat.  
    if (state[i] != eating)  
        self[i].wait();  
}
```

```
void putdown(int i) {  
    state[i] = thinking;  
    // test left and right neighbors  
    test((i+4) % 5);  
    test((i+1) % 5);  
}
```

# Dining Philosophers

```
void test(int i) {  
    if ( (state[(i + 4) % 5] != eating) &&  
        (state[i] == hungry) &&  
        (state[(i + 1) % 5] != eating)) {  
        state[i] = eating;  
        self[i].signal();  
    }  
}
```

# Dining Philosophers

```
dp.pickup[i];  
....  
eat  
....  
dp.putdown(i);
```

- It is easy to show that no two neighbors are eating simultaneously and no deadlocks will occur.
- However, it is possible for a philosopher to starve to death.

# Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next-count = 0;
// number of processes suspended on next.
```

- Each external procedure  $F$  will be replaced by **wait(mutex);**

...

body of  $F$ ;

...

```
if (next-count > 0)
  signal(next)
else
  signal(mutex);
```

- Mutual exclusion within a monitor is ensured.

# Monitor Implementation

- For each condition variable **x**, we have:

```
semaphore x-sem; // (initially = 0)
int x-count = 0;
```

- The operation **x.wait** can be implemented as:

```
x-count++;
if (next-count > 0)
    signal(next);
else
    signal(mutex);
wait(x-sem);
x-count--;
```

# Monitor Implementation

- The operation **x.signal** can be implemented as:

```
if (x-count > 0) {  
    next-count++;  
    signal(x-sem);  
    wait(next);  
    next-count--;  
}
```

# Monitor Implementation

## ■ *Conditional-wait* construct: **x.wait(c);**

- **c** – integer expression evaluated when the **wait** operation is executed.
- value of **c** (a *priority number*) stored with the name of the process that is suspended.
- when **x.signal** is executed, process with smallest associated priority number is resumed next.

## ■ Check two conditions to establish correctness of system:

- User processes must always make their calls on the monitor in a correct sequence.
- Must ensure that an uncooperative process does not ignore the mutual-exclusion gateway provided by the monitor, and try to access the shared resource directly, without using the access protocols.

# Solaris 2 Synchronization

- Solaris2 OS is designed to provide real-time computing, be multithreaded and support multiple processors.
- To control access to critical regions Solaris 2 implements
  - Adaptive mutexes
  - Condition variables
  - Semaphores
  - Reader-writer locks
  - Turnstiles
- Adaptive mutex protects access to every critical data item.
  - On multiprocessor system an adaptive mutex starts as a standard semaphore implemented as a spinlock.
  - Adaptive mutex is used to protect only data that are accessed by short-code segments (few hundred instructions).
- For longer code segments, condition variables and semaphores are used.
- Reader-writers locks are used to access data that is accessed frequently in read-only manner.
  - Semaphores serialize the access
  - When there are many readers and few writers r-w locks are efficient.
- Solaris 2 uses turnstiles to order list of threads waiting to acquire either an adaptive mutex or a reader-writer lock.



# Solaris 2 Synchronization

- Uses *turnstiles* to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock.
  - It is a queue structure containing threads blocked in a lock.
- To prevent a priority inversion, turnstiles are organized into priority inheritance protocol.
  - When a lower priority thread holds a lock, and higher priority thread blocks, the lower priority thread inherits the priority of the higher-priority thread.

# Windows 2000 Synchronization

- Multi-threaded kernel
  - Real-time applications and multiple processors
- Uses interrupt masks to protect access to global resources on uni-processor systems.
- Uses *spinlocks* on multiprocessor systems.
  - Kernel ensures that a thread will never be preempted while holding a spinlock.
- Also provides *dispatcher objects* which may act as mutexes and semaphores and events.
- Dispatcher objects may also provide *events*. An event acts much like a condition variable which may notify a waiting thread when a desired condition occurs.

# Transactional Memory

- Multi-core applications increases risk of race conditions and deadlocks.
- Techniques proposed: locks, semaphores and monitors
- Transactional memory provides alternative strategy
- A memory transaction is a sequence of memory read and write transactions that are atomic. If all the operations are completed the memory transaction is committed. Otherwise, the operations must be aborted. Such a feature can be added to programming language.
- Traditional way
  - ```
Update () {  
    Acquire();  
    /* Modify shared data */  
    release(); }
```
- Suppose we add an atomic operation
  - ```
Update() {  
    Atomic {  
    }  
}
```

# Atomic transactions

- A group of statements should be executed as a logical unit.
- More than mutual exclusion
- Notion of transaction has been emerged
  - ACID properties
  - Atomicity, consistency, Isolation and durability
- Atomicity: all or nothing
- Consistency: one consistent state to another consistent state
- Isolation: Parallel execution is serial
- Durability: Changes are permanent
- Two-phase locking and log based recovery methods are followed.
- Will be studied in database systems course

# Transactional Memory

## ■ Advantage

- System is responsible for guaranteeing the atomicity.

## ■ Transactional memory systems can be implemented in either software or in hardware.

- STM: software transactional memory
  - ✓ appropriate code is inserted by compiler
- HTM: Hardware transactional memory
  - ✓ Uses cache coherency protocols to support transaction memory

# Critical Regions

- High-level synchronization construct
- A shared variable **v** of type **T**, is declared as:

**v: shared T**

- Variable **v** accessed only inside statement  
**region v when B do S**

where **B** is a boolean expression.

- While statement **S** is being executed, no other process can access variable **v**.
- The expression **B** is Boolean expression which governs the access to the critical region.

# Critical Regions

- Regions referring to the same shared variable exclude each other in time.
- When a process tries to execute the region statement, the Boolean expression  $B$  is evaluated. If  $B$  is true, statement  $S$  is executed. If it is false, the process is delayed until  $B$  becomes true and no other process is in the region associated with  $v$ .
- If the following two statements are executed concurrently, it will be equivalent to the serial execution “S1 followed by S2” or “S2 followed by S1”.
  - Region  $v$  when (true) S1;
  - Region  $v$  when (true) S2;
- CR construct guards against simple errors associated with the semaphore solution.

# Example – Bounded Buffer

- Shared data:

```
struct buffer {  
    int pool[n];  
    int count, in, out;  
}
```



# Bounded Buffer Producer Process

- Producer process inserts **nextp** into the shared buffer

```
region buffer when( count < n) {  
    pool[in] = nextp;  
    in:= (in+1) % n;  
    count++;  
}
```

# Bounded Buffer Consumer Process

- Consumer process removes an item from the shared buffer and puts it in **nextc**

```
region buffer when (count > 0) {  
  nextc = pool[out];  
  out = (out+1) % n;  
  count--;  
}
```

# Implementation region $x$ when $B$ do $S$

- Can be implemented using semaphores.

# Implementing *S* as a Binary Semaphore

- Can implement a counting semaphore *S* as a binary semaphore.
- Data structures:  
**binary-semaphore S1, S2;**  
**int C;**
- Initialization:  
**S1 = 1**  
**S2 = 0**  
**C = initial value of semaphore S**

```
□ wait operation
  wait(S1);
  C--;
  if (C < 0) {
    signal(S1);
    wait(S2);
  }
  signal(S1);
```

```
□ signal operation
  wait(S1);
  C ++;
  if (C <= 0)
    signal(S2);
  else
    signal(S1);
```

## *Counting semaphores*

*wait(S):*

```
  S.value--;
  if (S.value < 0) {
    add this process to
    S.L;
    block;
  }
```

*signal(S):*

```
  S.value++;
  if (S.value <= 0) {
    remove a process P
    from S.L;
    wakeup(P);
  }
```